

# Migrating Amazon RDS for MySQL to Azure Database for MySQL using the MySQL Workbench Migration Wizard

You can use a variety of utilities, such as MySQL Workbench Export/Import, Azure Database Migration Service (DMS), and MySQL dump and restore, to migrate Amazon RDS for MySQL to Azure Database for MySQL. However, using the MySQL Workbench Migration Wizard provides an easy and convenient way to move your Amazon RDS for MySQL databases to Azure Database for MySQL.

With the Migration Wizard, you can conveniently select which schemas and objects to migrate. It also provides the ability to view server logs to identify errors and bottlenecks in real time. As a result, you can edit and modify tables or database structures and objects during the migration process when an error is detected, and then resume migration without having to restart from scratch.

**Note:** You can also use the Migration Wizard to migrate other sources, such as Microsoft SQL Server, Oracle, PostgreSQL, MariaDB, etc., but we'll cover those scenarios in separate blog posts.

## Prerequisites

Before you start the migration process, it is recommended that you ensure that several parameters and features are configured and set up properly, as described below.

- Make sure the Character set of the source and target databases are the same.
- Set the wait timeout to a reasonable time depending on the amount data or workload you want to import or migrate.
- Set the `max_allowed_packet` parameter to a reasonable amount depending on the size of the database you want to import or migrate.
- Verify that all of your tables use InnoDB, as Azure Database for MySQL Server supports only InnoDB Storage engine.
- Remove, replace, or modify all triggers, stored procedures, and other functions containing root user or super user definers (Azure Database for MySQL doesn't support the Super user privilege).

To replace the definers with the name of the admin user that is running the import process, run the following command:

```
DELIMITER; ;/*!50003 CREATE*/ /*!50017 DEFINER=`root`@`127.0.0.1`*/ /*!50003
DELIMITER;
/* Modified to */
DELIMITER;
/*!50003 CREATE*//*!50017 DEFINER=`AdminUserName`@`ServerName`*/ /*!50003
DELIMITER;
```

- If User Defined Functions (UDFs) are running on your database server, you need to delete the privilege for the mysql database. To determine if any UDFs are running on your server, use the following query:

```
SELECT * FROM mysql.func;
```

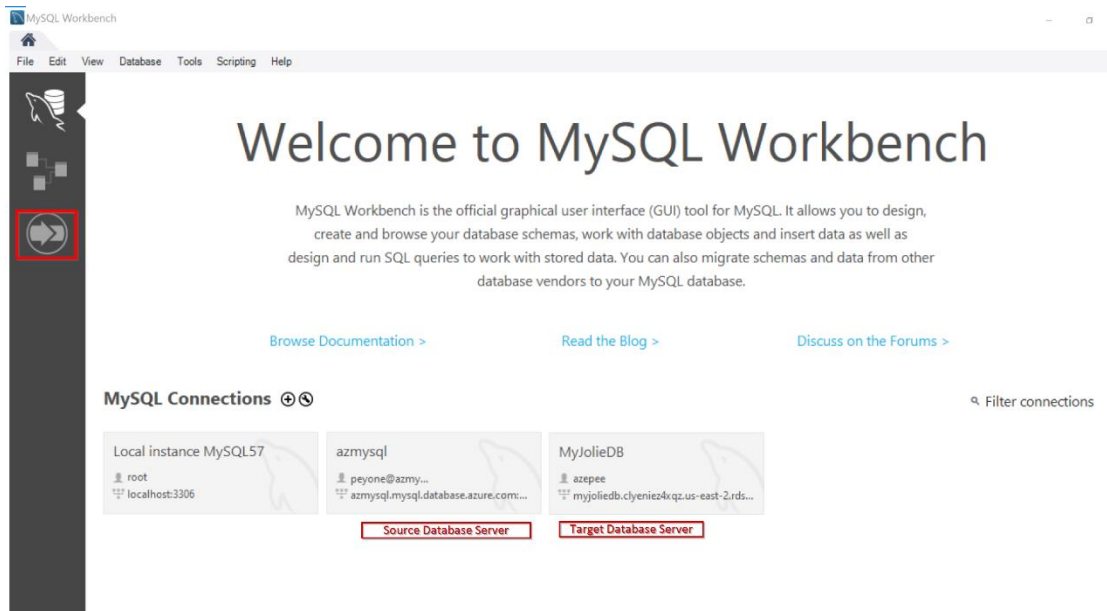
If you discover that UDFs are running, you can drop the UDFs by using the following query:

```
DROP FUNCTION your_UDFunction;
```

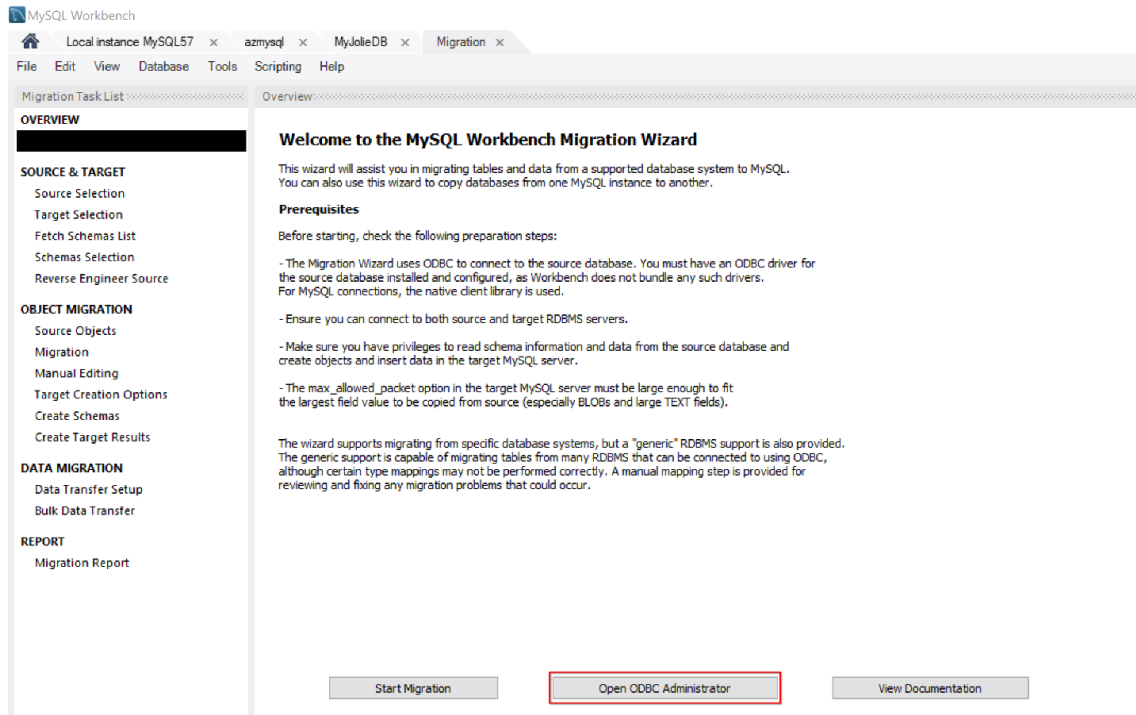
- Make sure that the server on which the tool is running, and ultimately the export location, has ample disk space and compute power (Vcores, CPU, and Memory) to perform the export operation, especially when exporting a very large database.
- Create a path between the on-premises or AWS instance and Azure Database for MySQL if the workload is behind firewalls or other network security layers.

## Begin the migration process.

1. To start the migration process, log onto MySQL Workbench, and then select the home icon.
2. In the left-hand navigation bar, select the Migration Wizard icon, as shown in the screenshot below.

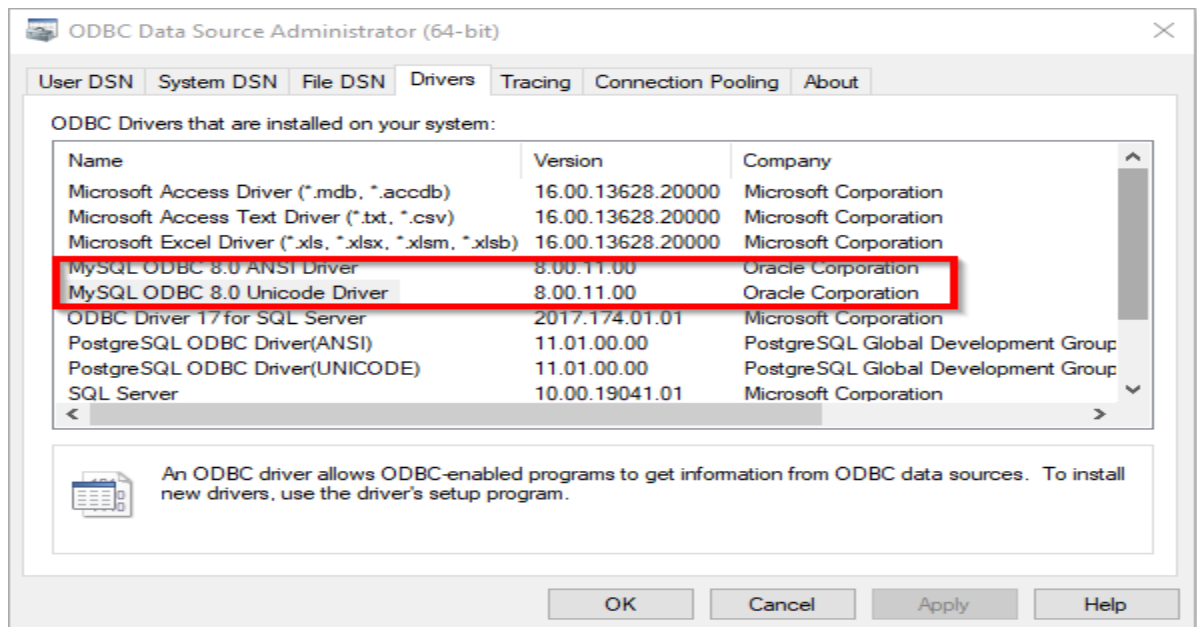


The **Overview** page of the Migration Wizard is displayed, as shown below.



3. Determine if you have an ODBC driver for MySQL Server installed by selecting **Open ODBC Administrator**.

In our case, on the **Drivers** tab, you'll notice that there are already two MySQL Server ODBC drivers installed.

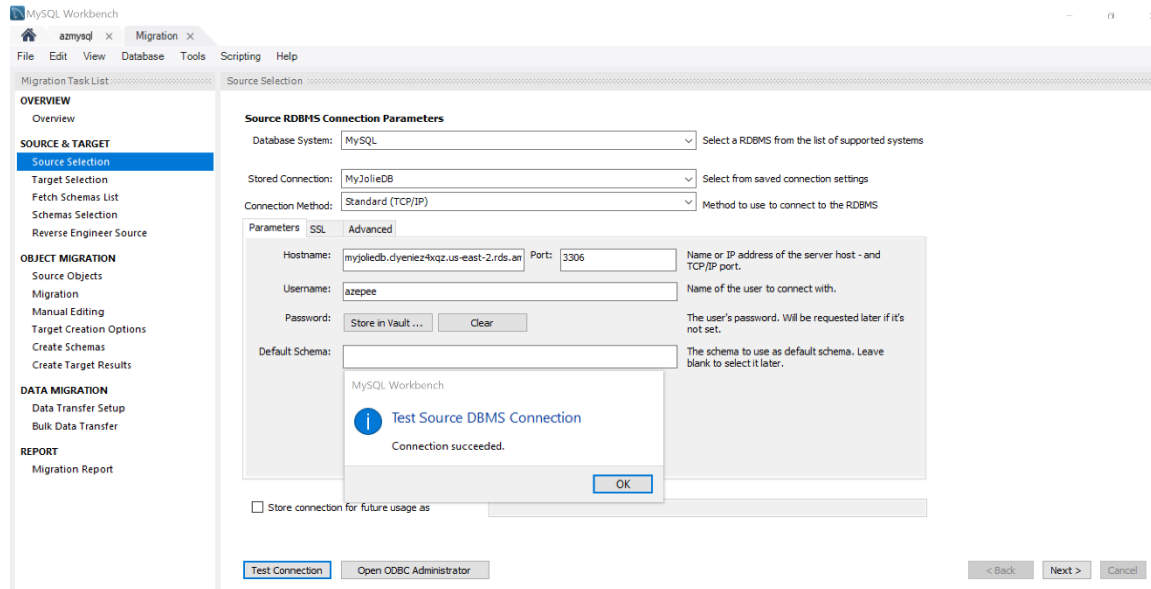


If a MySQL ODBC driver is not installed, use the MySQL Installer you used to install MySQL Workbench to install the driver. For more information about MySQL ODBC driver installation, see the following resources:

- [MySQL :: MySQL Connector/ODBC Developer Guide :: 4.1 Installing Connector/ODBC on Windows](#)
  - [ODBC Driver for MySQL: How to Install and Set up Connection \(Step-by-step\) – {coding}Sight \(codingsight.com\)](#)
4. Close the **ODBC Data Source Administrator** dialog box, and then continue with the migration process.

## Configuring connection parameters for the source database server

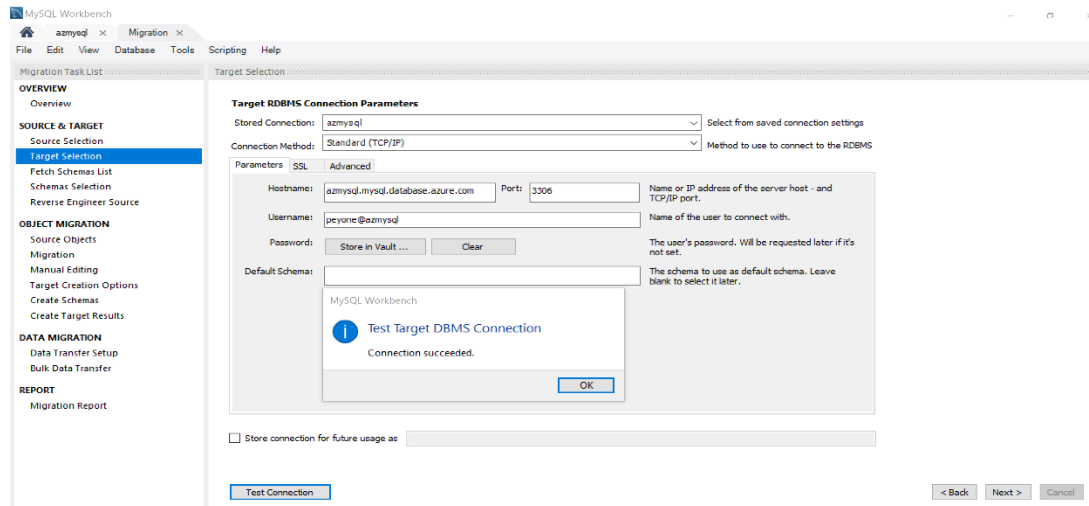
1. On the **Overview** page, select **Start Migration**.  
The **Source Selection** page appears. Use this page to provide information about the RDBMS you are migrating from and the parameters for the connection.
2. In the **Database System** field, select **MySQL**.
3. In the **Stored Connection** field, select one of the saved connection settings for that RDBMS.  
You can save connections by marking the checkbox at the bottom of the page and providing a name of your preference.
4. In the **Connection Method** field, select **Standard TCP/IP**.
5. In the **Hostname** text field, type the name of your source database server.
6. In the Port field, specify **3306**, and then enter the username and password to use to connect to the server.
7. In the **Database** field, enter the name of the database you want to migrate if you know it; otherwise leave this field blank.



8. Select **Test Connection** to check the connection to your MySQL Server instance.  
If you've entered the correct parameters, a message appears indicating a successful connection attempt.
9. Select **Next**.

## Configuring connection parameters for the target database server

1. On the **Target Selection** page, set the parameters to connect to your Destination MySQL Server instance using a process similar to that for setting up the connection to the source server.
2. When you are done, select **Test Connection** to verify a successful connection.



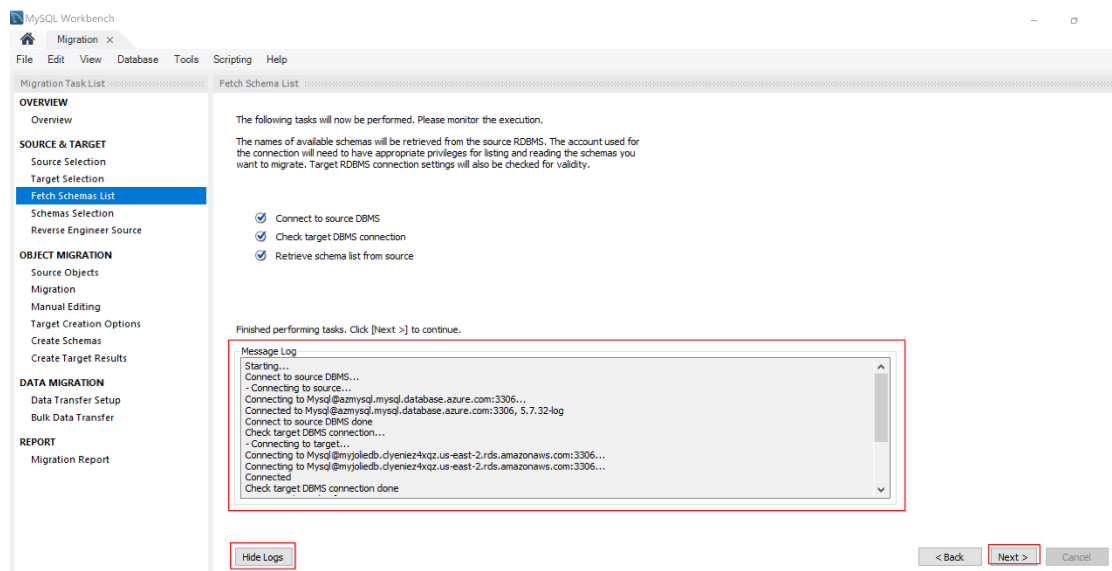
3. Select **Next**.

## Selecting the schemata to migrate

The Migration Wizard will communicate to your MySQL Server instance and fetch a list of schemata from the source server.

1. Select **show logs** to view this operation.

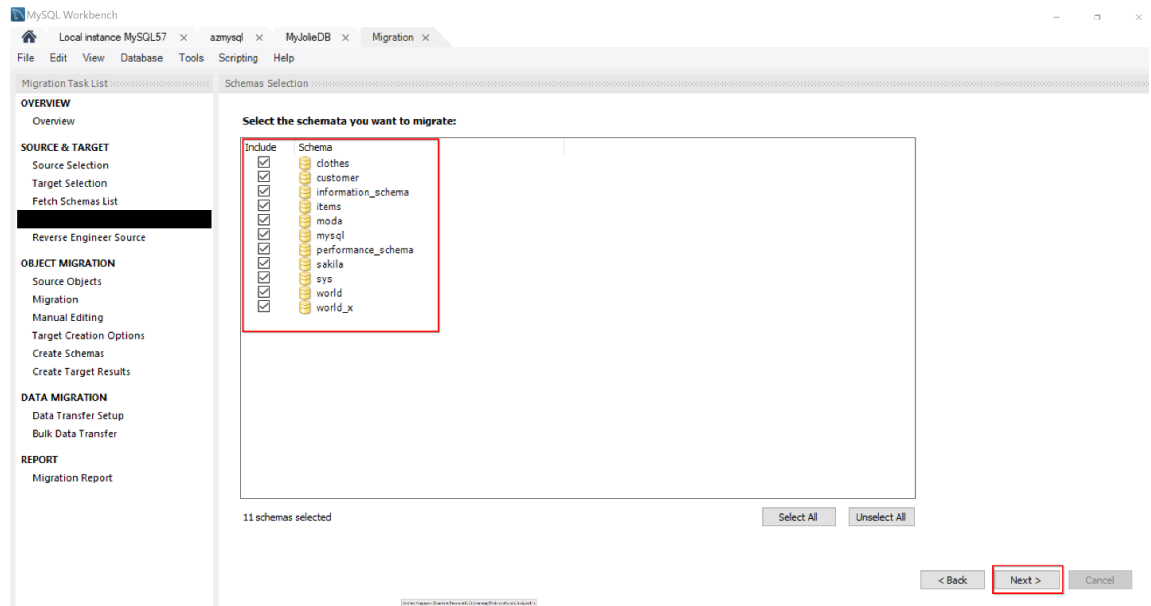
The screenshot below shows how the schemata are being retrieved from the source database server.



2. Select **Next** to verify that all the schemata were successfully fetched.

The screenshot below shows the list of fetched schemata.

3. Select the schemata that you want to migrate.



Keep in mind that you can only migrate schemata that appear in the list of schemata.

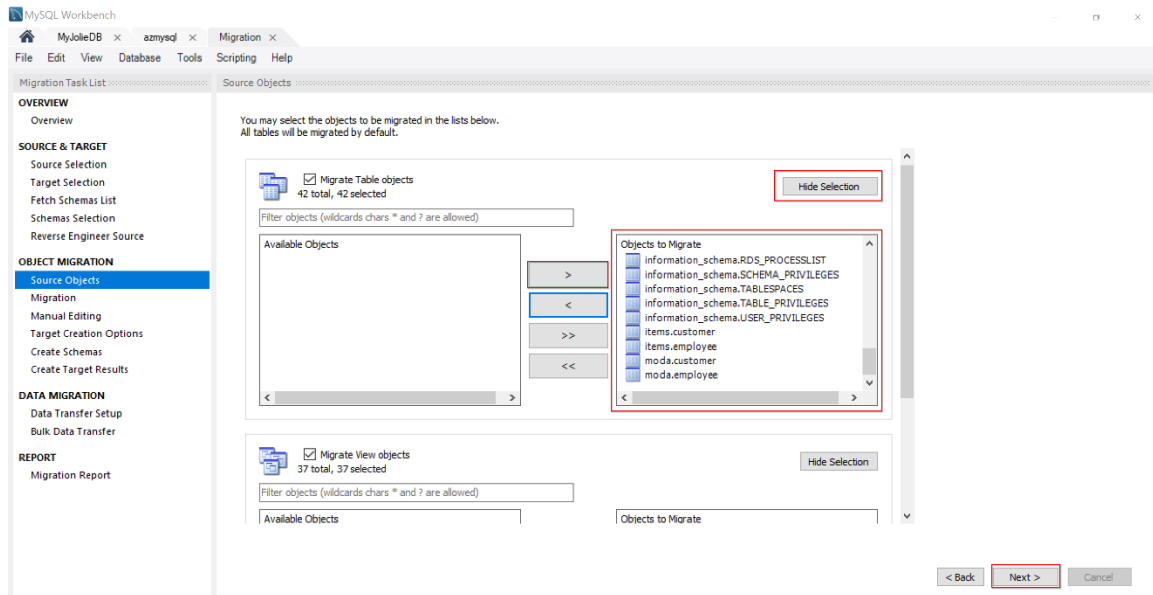
4. Select **Next**.

## Object migration

Next, specify the object(s) that you want to migrate.

1. Select **Show Selection**, and then, under **Available Objects**, select and add the objects that you want to migrate.

When you have added the objects, they will appear under **Objects to Migrate**, as shown in the screenshot below.



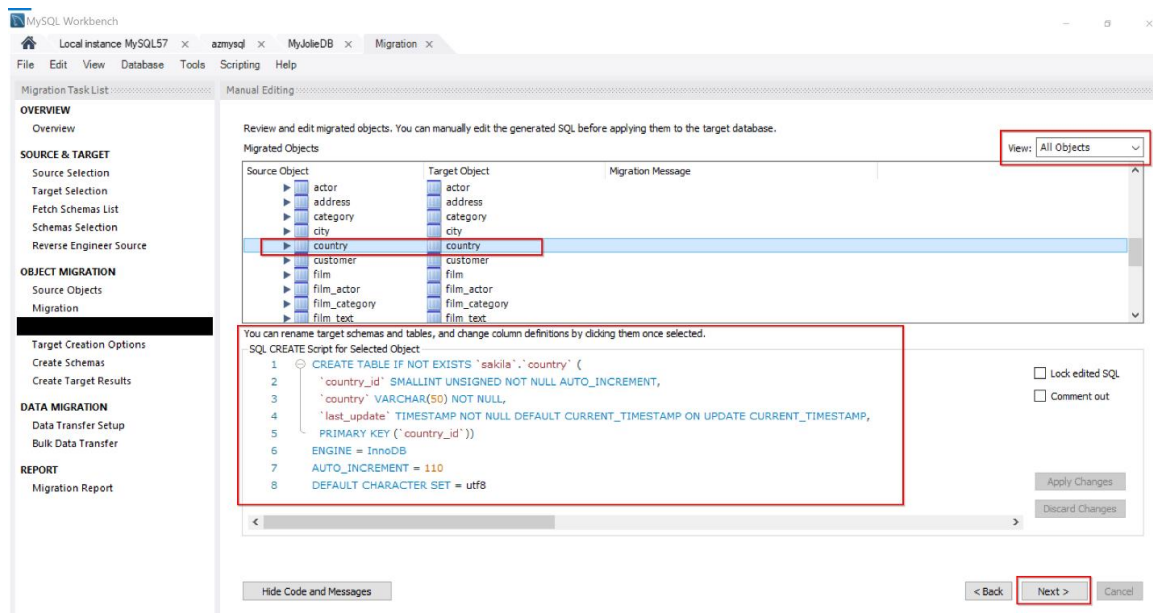
In this scenario, we've selected all table objects.

2. Select **Next**.

## Editing data

In this section you have the option of editing the objects that you want to migrate.

1. On the **Manual Editing** page, notice the **View** drop-down menu in the top-right corner.



The **View** drop-down box includes three items:

- **All Objects** – Displays all objects. With this option you can manually edit the generated SQL before applying them to the target database server. To do this, select the object and

select **Show Code and Messages**. You can see (and edit!) the generated MySQL code that corresponds to the selected object.

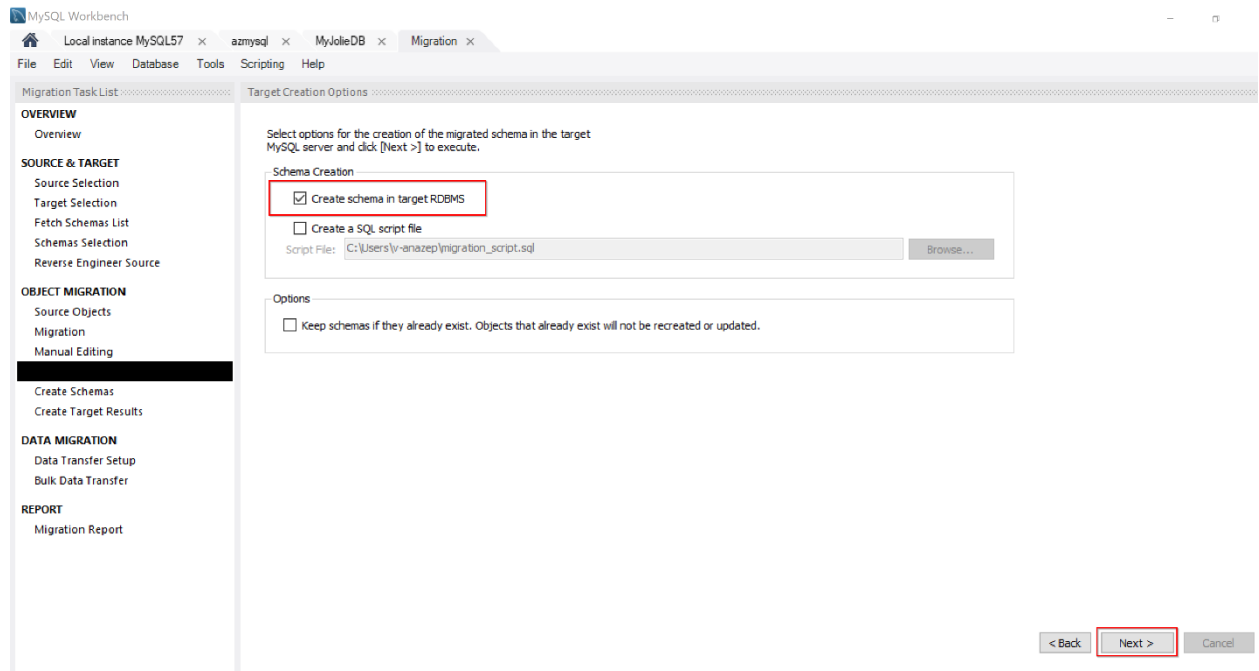
- **Migration problems** – Displays any problems that occurred during the migration, which you can review and verify.
- **Column Mapping** – Displays column mapping information. You can use this view to edit the name and change column of the target object.

2. Select **Next**.

## Creating the target database

1. Select the **Create schema in target RDBMS** check box.

You can also choose to keep already existing schemata, so they will not be modified or updated.



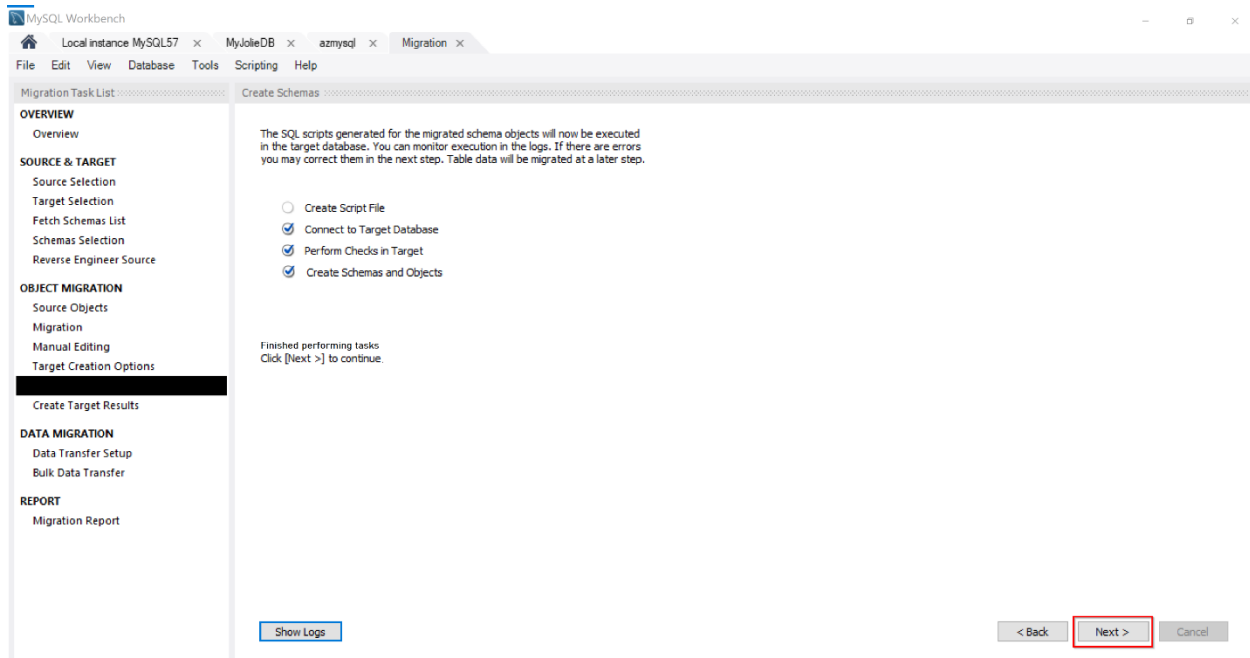
We'll choose to create schema in target RDBMS option for this blog, but you can still select both options, the other option being to create a SQL script file, so that you can save the script to a file on your local machine for other purposes.

2. Select **Next**.

## Running the MySQL script to create the database objects.

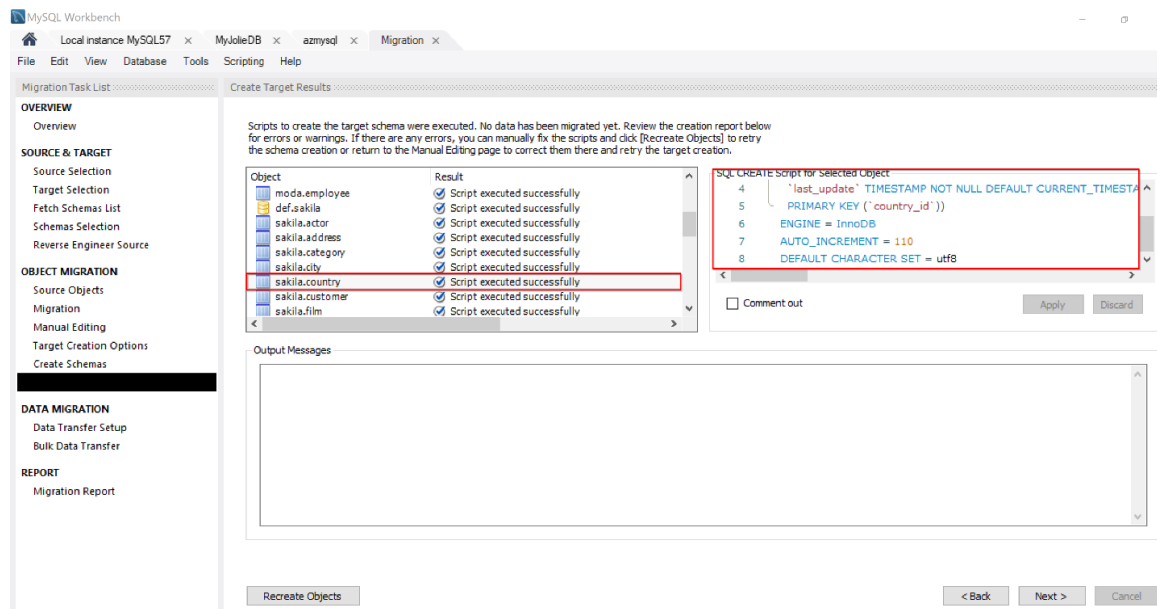
Since we have chosen the create schema in Target RDBMS option, the migrated SQL script will be executed in the target MySQL server. You can view its progress as shown in the screenshot below:





1. After creation of the schemata and their objects completes, select **Next**.

On the **Create Target Results** page, you're presented with a list of the objects created and notification of any errors that were encountered while creating them, as shown in the following screenshot.



2. Review the detail on this page to verify that everything completed as intended.

In this scenario we don't have any errors. If you do need to address any error messages, you can edit the migration script.

3. In the **Object** box, select the object that you want to edit.

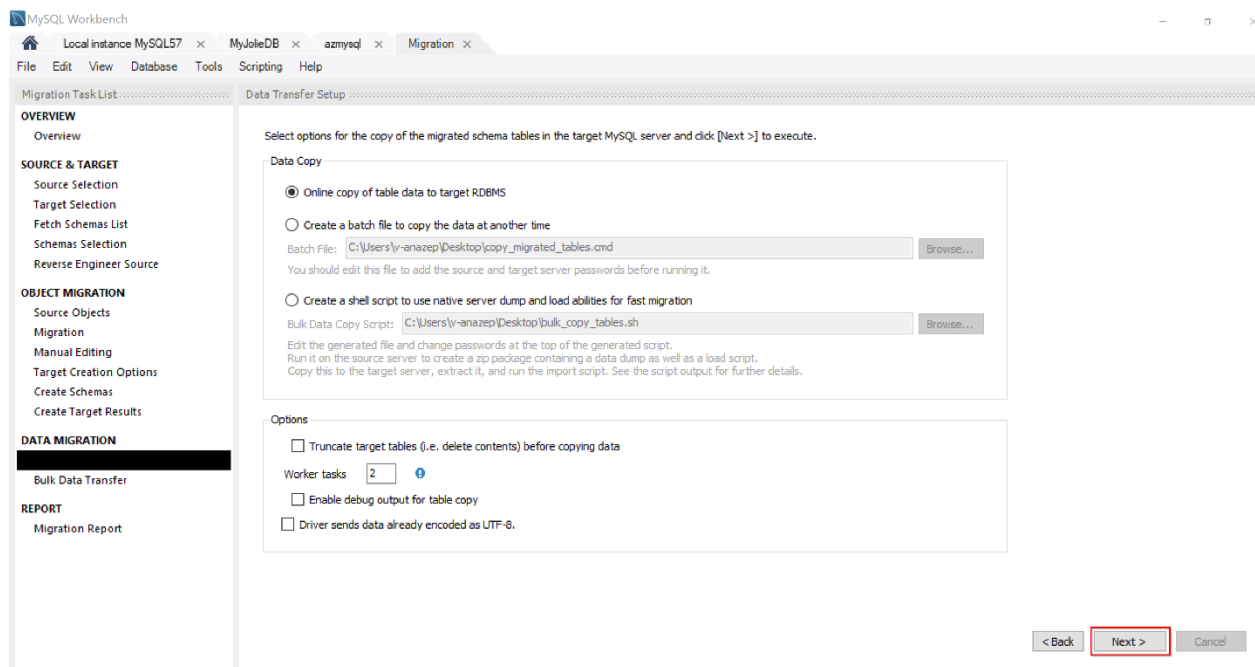
- Under **SQL CREATE script** for selected object, modify your SQL script, and then select **Apply** to save the changes.
- Select **Recreate Objects** to run the script including your changes.

If the script fails, you may need to edit the generated script. You can then manually fix the SQL script and run everything again. In this blog, we're not changing anything, so we'll leave the script as it is.

- Select **Next**.

## Transferring Data

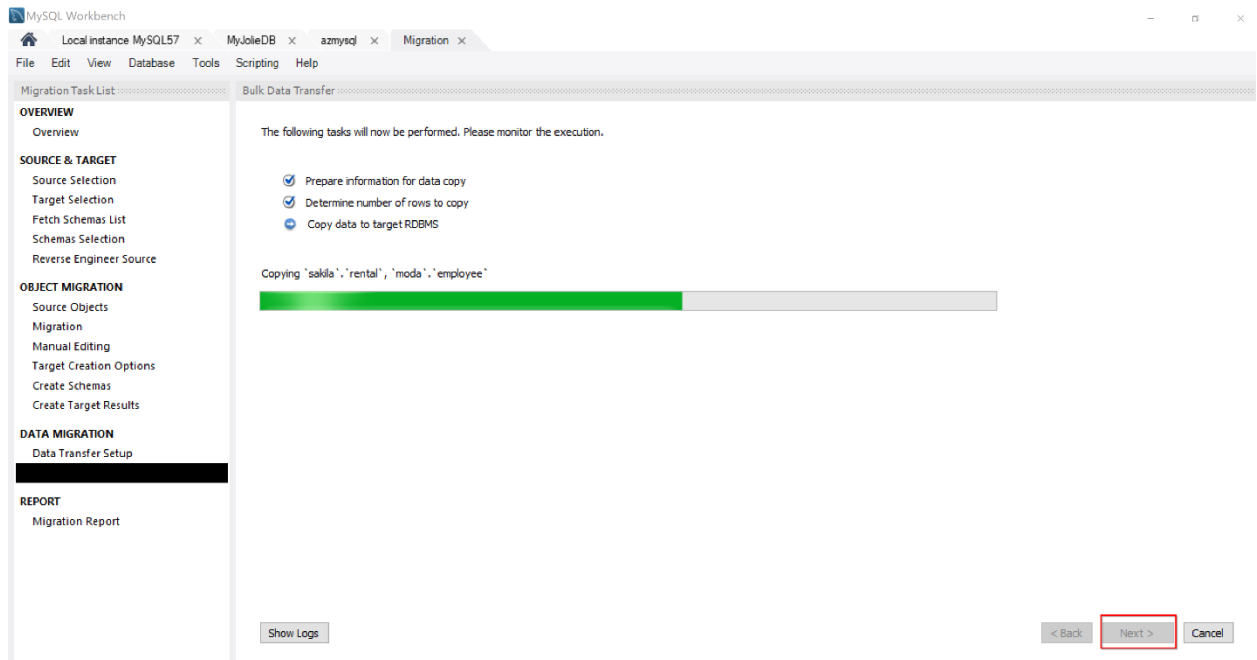
This part of the process moves data from the source MySQL Server database instance into your newly created target MySQL database instance. Use the **Data Transfer Setup** page to configure this process.



This page provides two options. One option allows you to perform a live/online transfer or to dump the data into a batch or a shell file that you can run later. The other option provides a way to tune this process. For the purposes of this post, we'll leave the default values for the options.

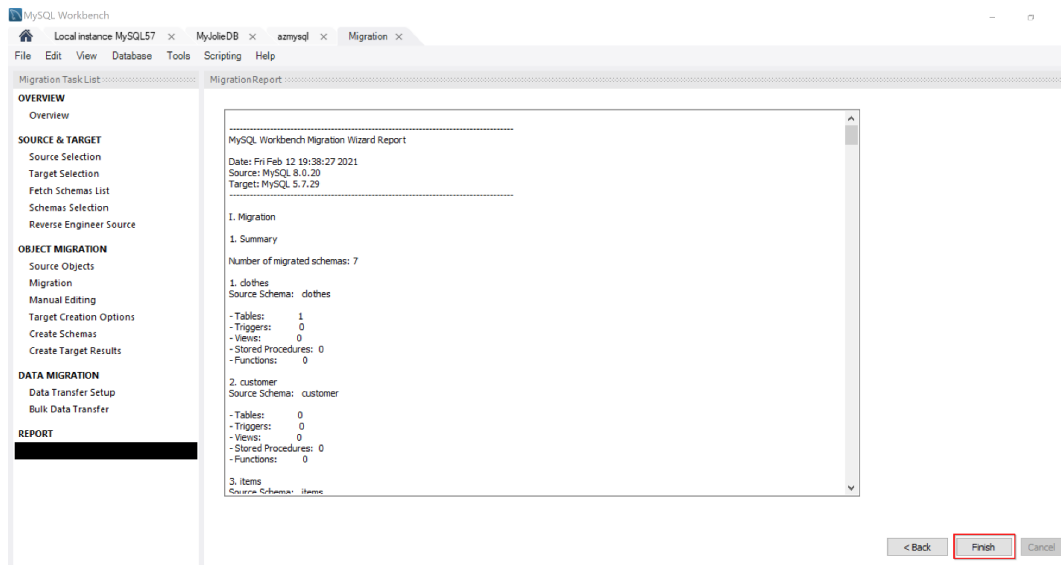
- To begin the actual process of transferring data, select **Next**.

It will take a little while to copy the data. The progress of the data transfer process appears as shown in the following screenshot.



2. After the transfer completes, select **Next**.

The **Migration Report** page appears, providing a report summarizing the whole process, as shown on the screenshot below:



3. Select **Finish** to close the Migration Wizard.

The migration is now successfully completed.

## Verifying consistency of the migrated schemata and tables

1. Next, log into your MySQL target database instance to verify that the migrated schemata and tables are consistent with your MySQL source database.

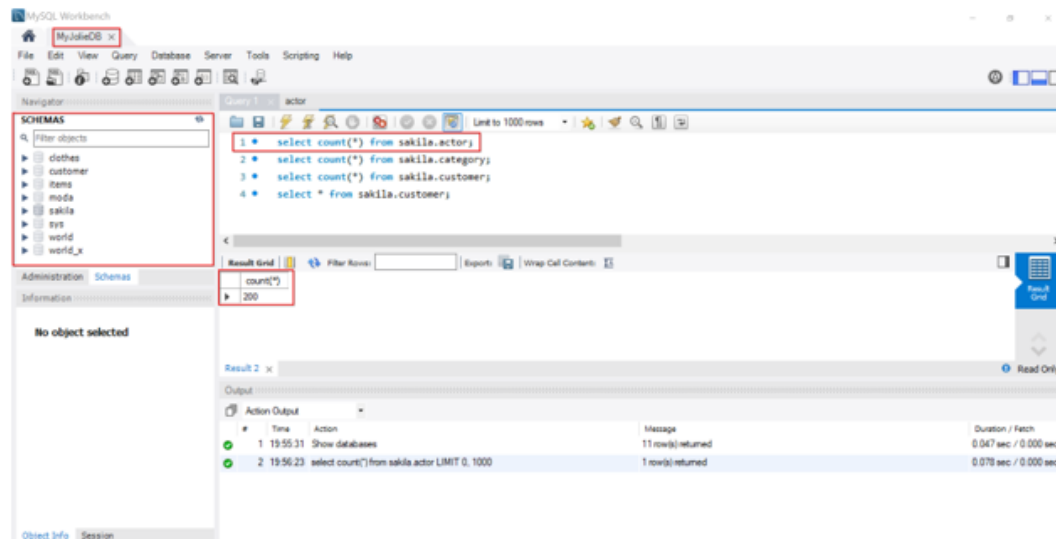
In our case, we can see that all schemata (sakila, moda, items, customer, clothes, world, and world\_x) from the Amazon RDS for MySQL: **MyjolieDB** database have been successfully migrated to the Azure Database for MySQL: **azmysql** instance.

2. To verify the table and rows counts, run the following query on both instances:

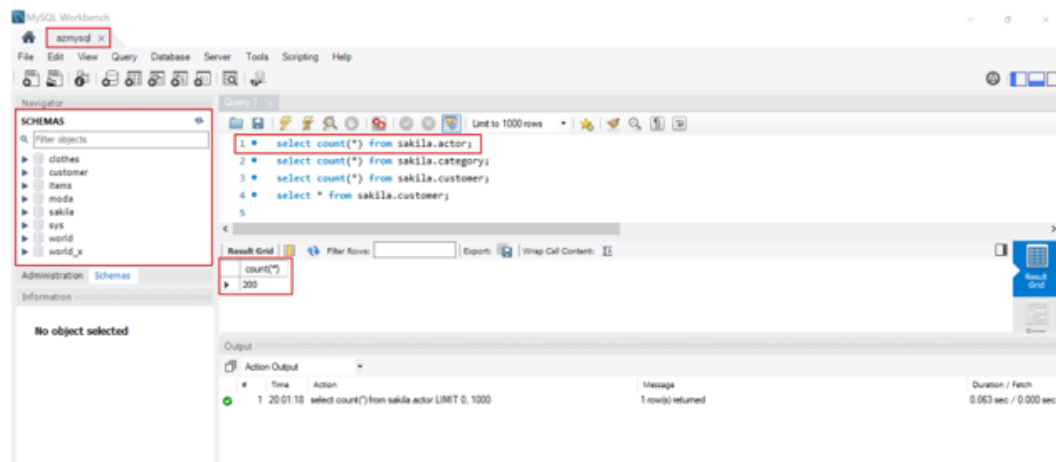
```
SELECT COUNT (*) FROM sakila.actor;
```

You can see from the screenshot below that the row count for Amazon RDS MySQL is 200, which matches the Azure Database for MySQL instance.

AWS\_RDS MySQL Instance: myjoliedb.clyniez4xqz.us-east-2.rds.amazonaws.com



Azure Database for MySQL Instance: azmysql.mysql.database.azure.com



While you can run the above query on every single schema and table, that will be quite a bit of work if you're dealing with hundreds of thousands or even millions of tables. You can use the queries below to verify the schema (database) and table size instead.

3. To check the database size, run the following query:

```
SELECT table_schema AS "Database",
```

```
ROUND(SUM(data_length + index_length) / 1024 / 1024, 2) AS "Size (MB)"
FROM information_schema.TABLES
GROUP BY table_schema;
```

4. To check the table size, run the following query:

```
SELECT table_name AS "Table",
ROUND(((data_length + index_length) / 1024 / 1024), 2) AS "Size (MB)"
FROM information_schema.TABLES
WHERE table_schema = "database_name"
ORDER BY (data_length + index_length) DESC;
```

You see from the screenshot below that schema (database) size from the Source Amazon RDS MySQL instance is the same as that of the target Azure Database for MySQL Instance.

AWS\_RDS MySQL Instance: myjoliedb.clyniez4xqz.us-east-2.rds.amazonaws.com

The screenshot shows the MySQL Workbench interface for an AWS RDS instance named 'MyJolieDB'. The 'SCHEMAS' panel on the left lists various databases. The 'Query' editor in the center contains a SQL query to calculate the size of each schema. The 'Result Grid' at the bottom displays the results of this query.

Database	Size (MB)
clothes	0.02
customer	0.00
items	0.03
mode	0.03
mysql	7.73
performance_schema	0.00
sakila	6.45
sys	0.02
world	0.77
world_x	0.84

Azure Database for MySQL Instance: azmysql.mysql.database.azure.com

The screenshot shows the MySQL Workbench interface for an Azure Database for MySQL instance named 'azmysql'. The 'SCHEMAS' panel on the left lists various databases. The 'Query' editor in the center contains the same SQL query as the previous screenshot. The 'Result Grid' at the bottom displays the results of this query.

Database	Size (MB)
clothes	0.02
customer	0.00
items	0.03
mode	0.03
mysql	7.73
performance_schema	0.00
sakila	6.45
sys	0.02
world	0.77
world_x	0.84

Since the schema (database) sizes are the same in both instances, you don't really need to check individual table sizes. Nevertheless, you can always use the above query to check your table sizes, as necessary.

You've now confirmed that your migration completed successfully.

## Conclusion

This concludes my post on migrating Amazon RDS for MySQL to Azure Database for MySQL using the MySQL Workbench Migration Wizard. I hope it's provided you with the information you need to perform this important task. If you have questions, please feel free to reach out to me or the Ask Azure DB for MySQL alias. Thank you!